

CMIP6 Amazon-hosted Data Informational Session

Aparna Radhakrishnan

Kristopher Rand

Charles Stern


Julius Busecke



CMIP6 Data Informational session, Oct 13th, 2021

Special thanks to..

Ryan Abernathey, V. Balaji, Philip Kershaw, Ana Privette, Ag Stephens, Naomi Naik, Serguei Nikonov, Hans Vahlenkamp, Mackenzie Blanus, Anderson Banihirwe, Chris Blanton, Nkeh Perry Boh, Ben Evans, Richard Smith, Rhys Evans, Zac Flamig, Diana Gergel, Thomas Jackson, Rebecca Monge, Natalie O'Leary, Zouberou Sayibou.

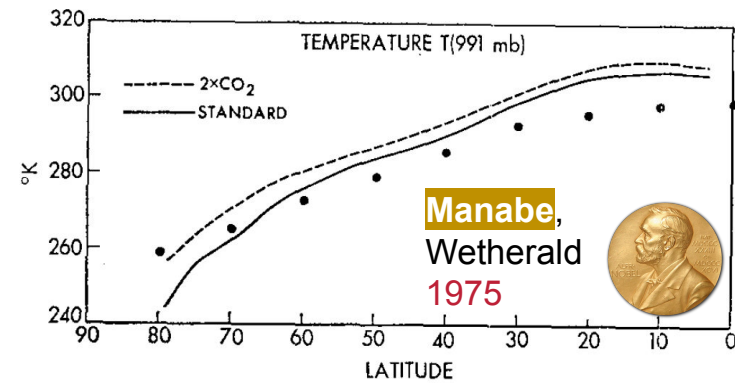


The Climate Pledge


Amazon is committed to building a sustainable business for our customers and the planet. In 2019, Amazon co-founded The Climate Pledge—a commitment to be net-zero carbon across our business by 2040, 10 years ahead of the Paris Agreement.



Motivation: Nature



3 (NOAA GFDL)+ 2 (NASA GISS) -- In Charney's report, 1979



IPCC AR6, Chapter 4, Fi 4.2 Global temperature change from CMIP6 historical and scenario simulations.

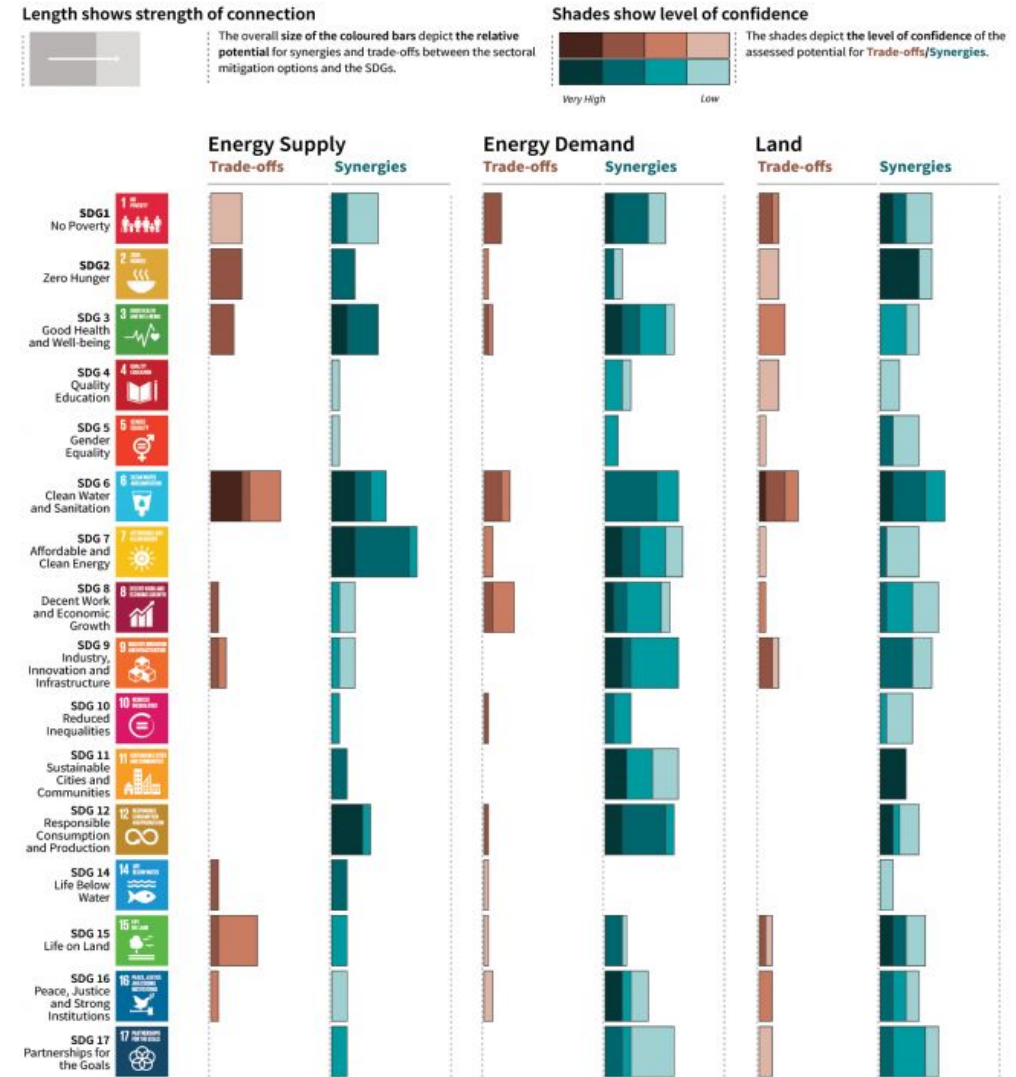


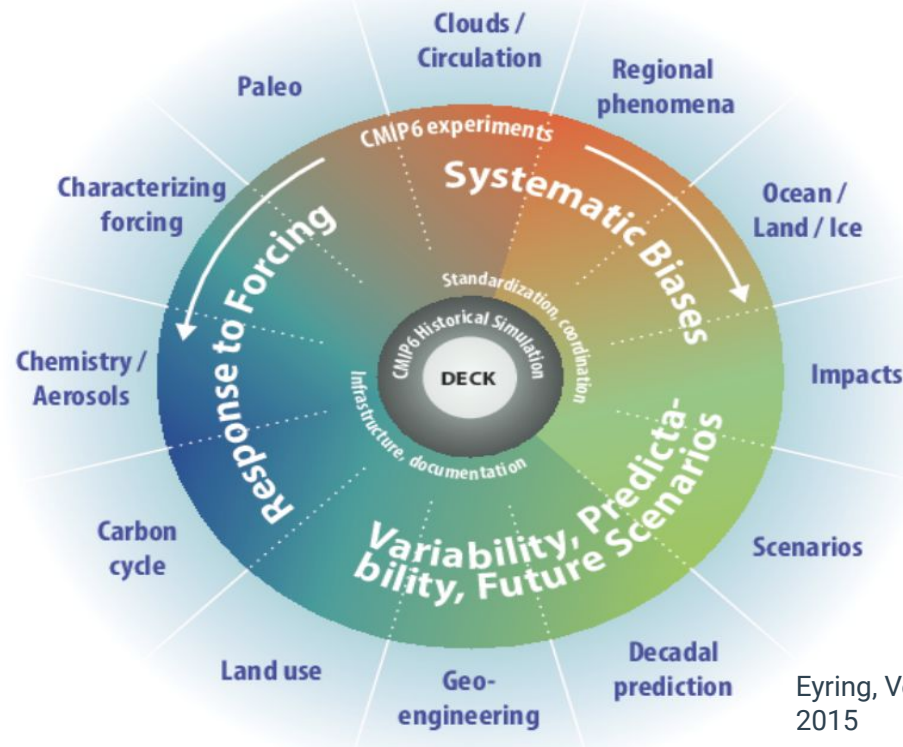

Photo: Saikiran Kesari/Unsplash

<https://science.thewire.in/environment/what-the-ipcc-report-forecasts-for-india-development-future/>



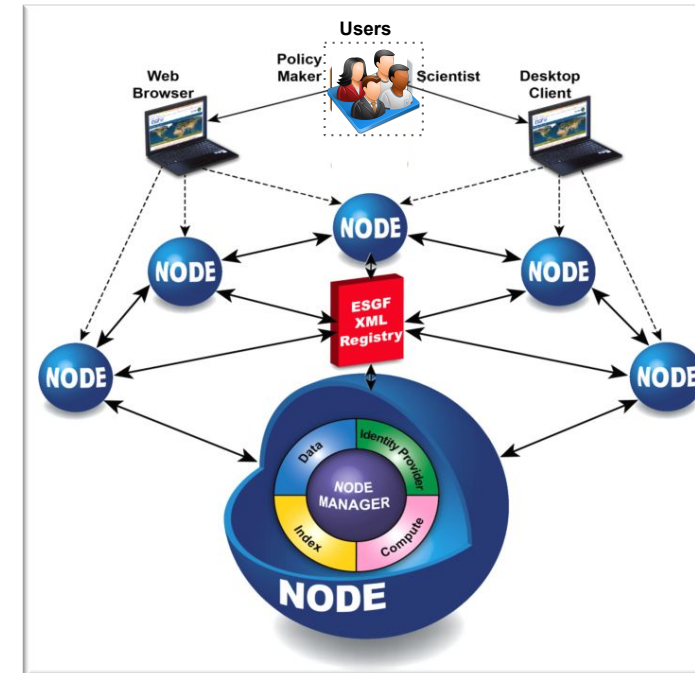
IPCC, 2021: Summary for Policymakers. In: Climate Change 2021: The Physical Science Basis. Contribution of Working Group I to the Sixth Assessment Report of the Intergovernmental Panel on Climate Change [MassonDelmotte, V., P. Zhai, et al.

Overview: ESGF and CMIP



Eyring, Veronika et al. 2015

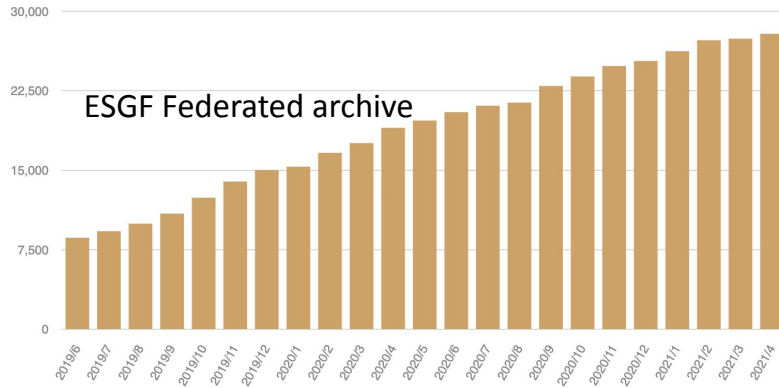
The objective of CMIP is to better understand past, present, and future climate change arising from natural, unforced variability or in response to changes in radiative forcings in a multi-model context.




Adapted from
<https://esgf.github.io/>

The Earth System Grid Federation (ESGF) is an international collaboration for the software that powers most global climate change research, notably assessments by the Intergovernmental Panel on Climate Change (IPCC).

Challenges and path forward




Expanding data archive..



Accelerate research




Different experimental design, model output formats, variable names, units, resolution,....

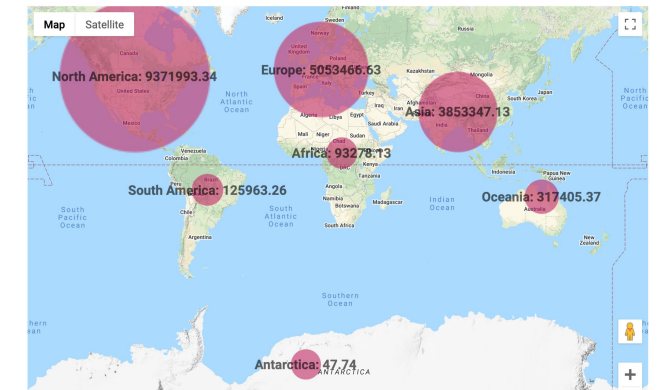


Democratize access to climate data

CMIP6 Data Informational Session, A.Radhakrishnan et al, Oct 13 2021



Numerous science questions..



Diverse groups of CMIP6 users all over the world. Cr. CMCC

Community-driven research efforts: CMIP6 data in the cloud



CMIP6 S3 bucket



arn:aws:s3:::esgf-world

[AWS CLI](#) Access aws s3 ls
s3://esgf-world/ --no-sign-request

<https://esgf-world.s3.amazonaws.com/index.html>

Intake-esm catalog

<https://cmip6-nc.s3.amazonaws.com/esgf-world.csv.gz>

THREDDS catalog

<https://aws-cloudnode.esgf.io/thredds/catalog/catalog.html>

SpatioTemporal Asset Catalogs (STAC) underway

arn:aws:s3:::cmip6-pds

[AWS CLI](#) Access aws s3 ls
s3://cmip6-pds --no-sign-request

<https://cmip6-pds.s3.amazonaws.com/index.html#CMIP6/>

Intake-esm catalog

<https://cmip6-pds.s3.amazonaws.com/pan-geo-cmip6.csv>

STAC catalogs underway

Checkout the [CMIP6 registry in AWS](#) to read more information, including CMIP6 data citations.

Community-driven best practices for improved data exploration

E.g. Directory Reference Syntax (DRS) established by the ESGF community makes cataloguing possible.

[Taylor et al.2017](#), [CMIP6-CV](#)

E.g.
`s3://esgf-world/CMIP6/AerChemMIP/NOAA-GFDL/GFDL-ESM4/hist-piNTCF/r1i1p1f1/Amon/tas/gr1/v20180701/tas_Amon_GFDL-ESM4_hist-piNTCF_r1i1p1f1_gr1_185001-194912.nc`



Intake-esm



Search for all atmos monthly surface temperature fields for historical simulations.

```
exp_filter = ['historical']
table_id_filter = 'Amon'
variable_id_filter = "tas"
cat = col.search(experiment_id=exp_filter,
                 table_id=table_id_filter,
                 variable_id=variable_id_filter)
```

catalog with 55 dataset(s) from 1872 asset(s):

intake-esm <https://intake-esm.readthedocs.io/en/latest/>

CMIP6 Controlled Vocabulary: https://github.com/WCRP-CMIP/CMIP6_CVs

intake-esm <https://intake-esm.readthedocs.io/en/latest/>

Example notebooks: <https://github.com/pangeo-data/pangeo-example-notebooks>

<https://github.com/aradhakrishnanGFDL/gfdl-aws-analysis/blob/master/examples/intake-esm-s3-nc-simple-access.ipynb>

Pangeo documentation: <https://pangeo-data.github.io/pangeo-cmip6-cloud/>


Community-driven research efforts: Bring analysis to data



ESGF federated nodes across the world

Towards a more accessible, discoverable and performant CMIP6 data holding in the cloud.

Earth System Grid Federation (ESGF) in the cloud. -By Kristopher Rand



<https://earthdata.nasa.gov/>

Towards a generic, flexible API to create analysis-ready cloud-optimized dataset. - By Charles Stern

Speed up your multi-model analysis, using CMIP6 pre-processor and Dask - By Julius Busecke

CMIP6 Data in the Cloud: 3 Main Objectives

Accessibility



Discoverability




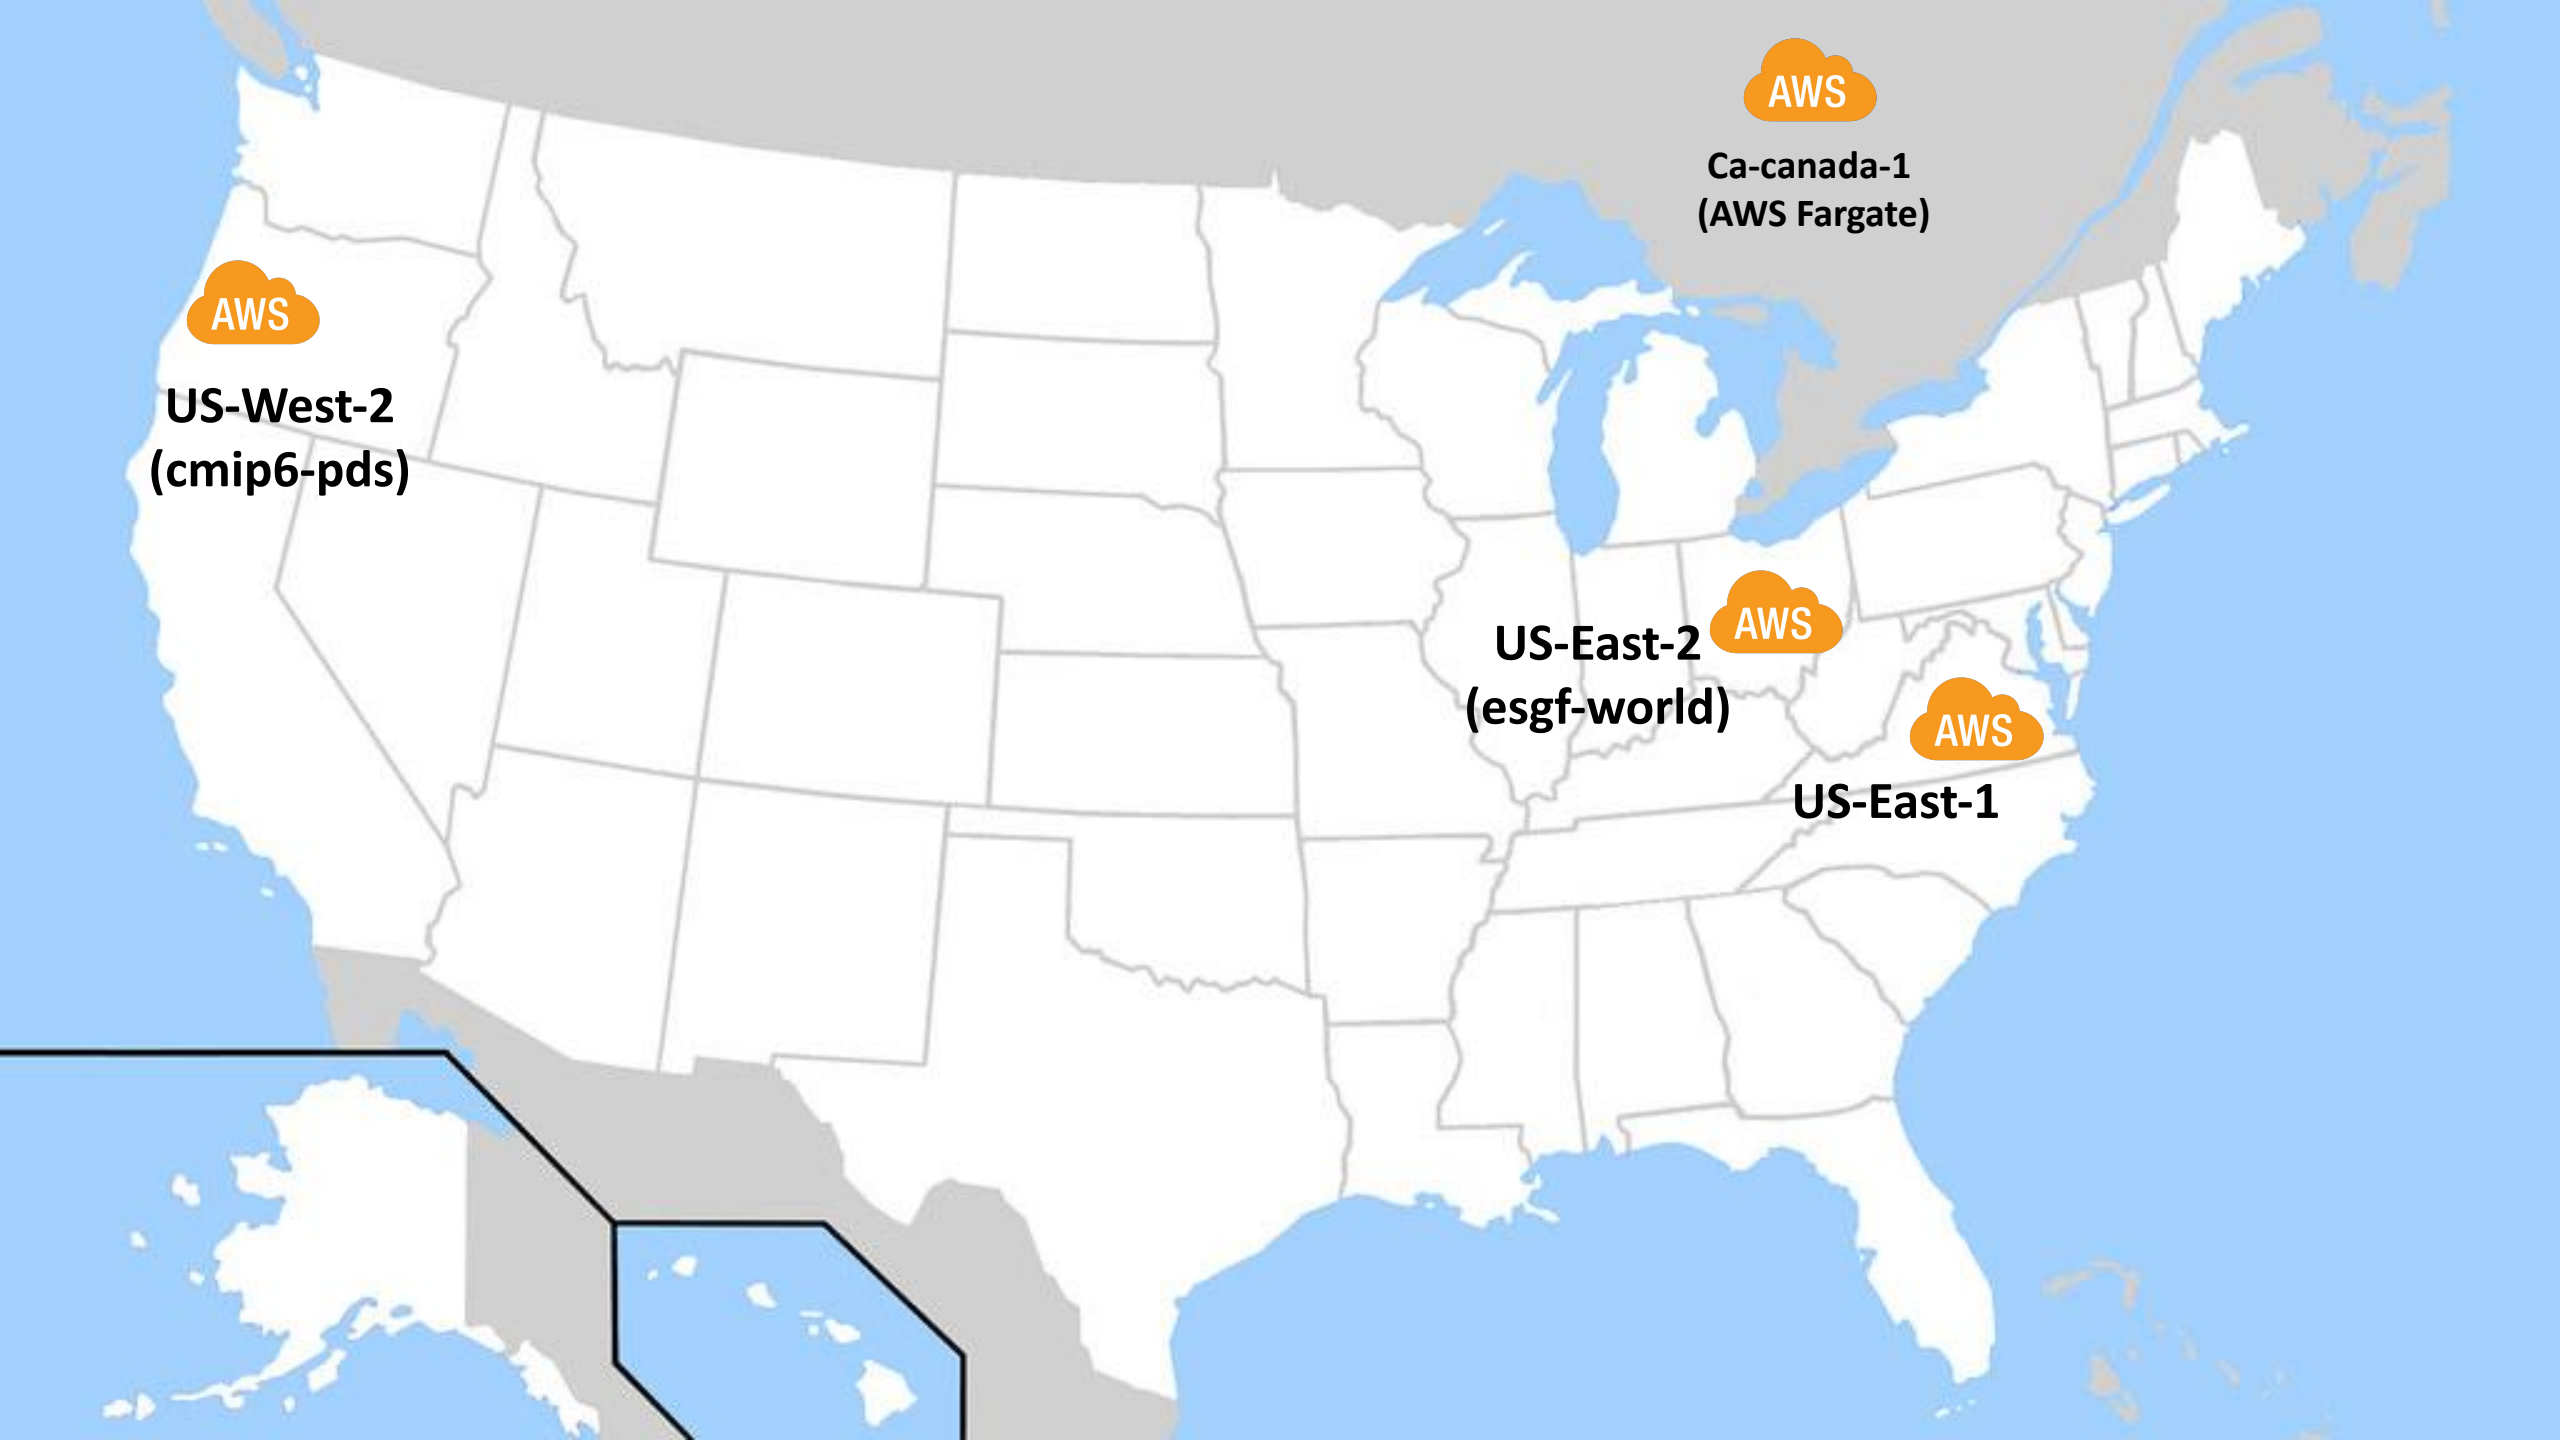
Performance



Others: 1. Scalable framework, 2. Data Management, 3. Data Metrics Analytics

Building a CMIP6 data node in the cloud utilizing Amazon's Elastic Kubernetes Service (EKS)





AWS

Ca-canada-1
(AWS Fargate)

AWS

US-West-2
(cmip6-pds)

AWS

US-East-2
(esgf-world)

AWS

US-East-1

EKS Clusters by Region



Ohio Region - US-East-2

- ESGF cluster
- 3 EC2 instances
- ESGF Data publication software
- Mounts “esgf-world” S3 bucket (1 PB) as file system



Virginia Region - US-East-1

- JupyterHub Dev/Test cluster
- 3-20 EC2 instances for autoscaling tests



Oregon Region - US-West-2

- ESGF Dev/Test cluster
- 1 EC2 instance
- The region itself also houses “cmip6-pds” Zarr S3 bucket (1 PB)



Canada Region - ca-canada-1

- AWS Fargate (serverless containers)

ESGF in the Cloud: Cloud Storage

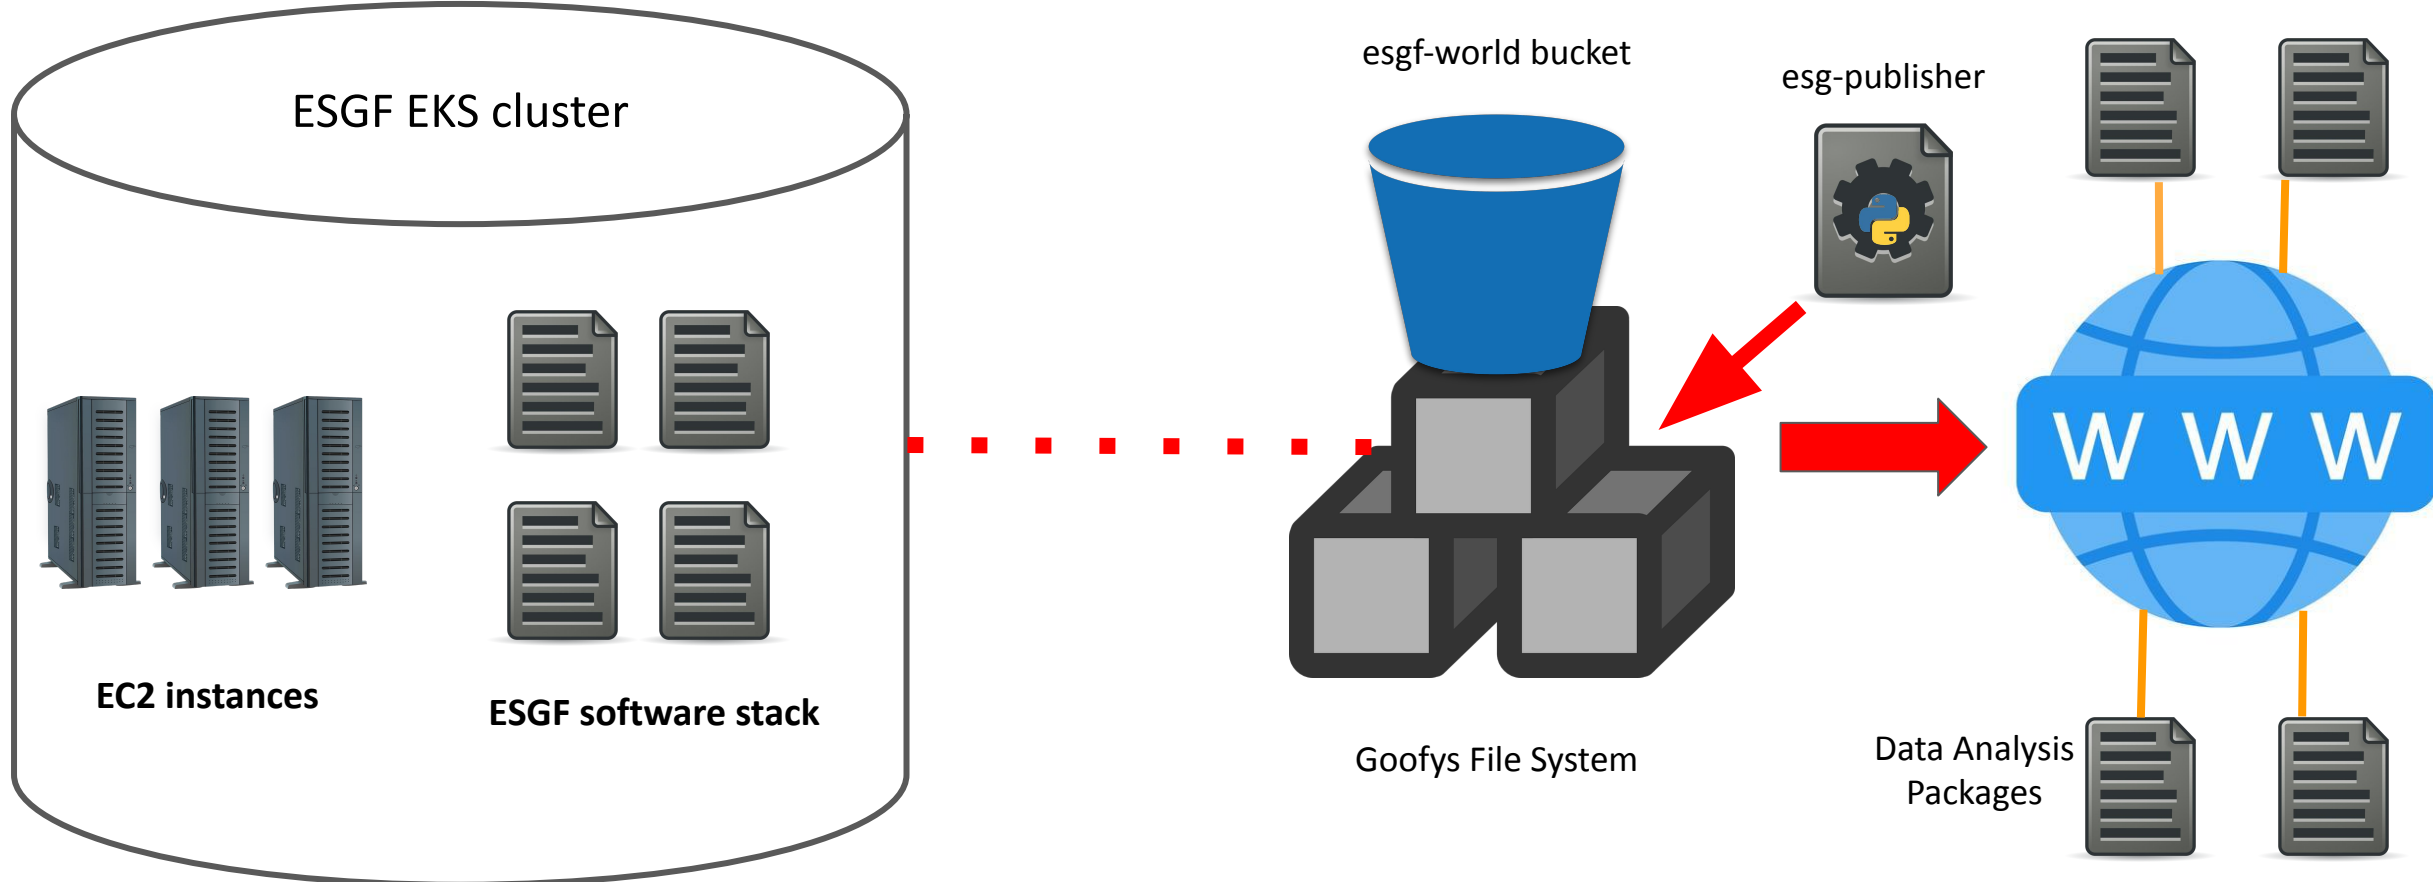
With our EKS clusters established, we wanted to anchor data holdings to a storage infrastructure that provided sufficient security, organization, and efficiency




- Amazon Simple Storage Service (S3) bucket
 - “esgf-world” bucket (NetCDF data) mounted to EC2 instances on ESGF cluster (1 PB storage)
 - “cmip6-pds” bucket (Zarr data) utilized on EC2 instance for ESGF dev/test cluster (1 PB storage)
- Dynamically optimized for performance and cost effectiveness depending upon user access
- Security policies: S3 bucket policies, object access and bucket access control lists (ACLs), IAM (Identify and Access Management) roles.
- Objects organized compliant to CMIP6 Data Reference Syntax (DRS)
- 2 PB of storage total between the 2 buckets

ESGF Cluster and the “esgf-world” S3 bucket

- EC2 instances/nodes mount S3 bucket using the POSIX-like file system, Goofys
 - **ESGF software stack** for eventual data publication in AWS requires interaction with the Amazon S3 bucket to perform R/W operations; cannot natively communicate with the S3 bucket
 - Goofys is high-performance and easy to set up



ESGF Cluster Architecture



Filling the bucket




Initial Data Transfer (esgf-world bucket)

- Ingress: Transfer of $\cong 650$ TB of high-value CMIP6 datasets primarily based on IPCC chapter variables and the consensus of the community
- Time period of transfer: 6 months




650 TB CMIP6 data reservoir



Transfer using AWS CLI


esgf-world bucket

Access from the esgf-world S3 bucket



THREDDS Data Server (TDS)

- Web server that provides metadata and data access for scientific datasets, using OPeNDAP, OGC WMS and WCS, HTTP, and other remote data access protocols.
- ESGF IO domain (<https://aws-cloudnode.esgf.io/thredds/catalog/catalog.html>)




ESGF THREDDS Data Server

Catalog

Dataset

- CMIP6
- AerChemMIP
- C4MIP
- CDRMIP
- CFMIP

A red arrow points from this screenshot to the one on the right.




ESGF THREDDS Data Server

Catalog

Dataset	Size	Last Modified
CMIP/NOAA-GFDL/GFDL-ESM4/1pctCO2/r1i1p1f1/Amon/ts/gr1/v20180701/		--
ts_Amon_GFDL-ESM4_1pctCO2_r1i1p1f1_gr1_000101-010012.nc	141.3 Mbytes	2020-11-14T02:29:49Z
ts_Amon_GFDL-ESM4_1pctCO2_r1i1p1f1_gr1_010101-015012.nc	70.48 Mbytes	2020-11-14T02:30:02Z

ESGF Data Publisher




- The ESGF data publisher is run on the datasets and is processed into a record that becomes visible to applications that utilize the ESGF search API (i.e. Synda search)
- Further enhances discovery capability by enabling bulk search parameters using dozens of available dataset metadata fields
- As the publisher is decoupled from the main ESGF software package, there is greater flexibility for data egress

Next Steps: Egress - Federation of ESGF Cloud node

aws-cloudnode.esgf.io data node



esgf-world S3 bucket



Home

WARNING: Not all models include a variant "r1i1p1f1", and across models, identical values of variant_label do not imp were used in each variant, please check modeling group publications and documentation provided through ES-DOC.

Enter Text:


Display 10

☐ Show All Replicas ☐ Show All Versions ☐ Search Local Node Only (Inc

The search returned 0 results.

MIP Era +
Activity +
Model Cohort +
Product +
Source ID +
Institution ID +
Source Type +
Nominal Resolution +
Experiment ID +
Sub-Experiment +
Variant Label +
Grid Label +
Table ID +
Frequency +
Realm +
Variable +
CF Standard Name +
Data Node +


Next Steps: Ingress - Filling our Bucket Faster



- Extending the ingress repository to the ESGF network itself
 - 11+ PB of unique datasets
 - Plentiful metadata
 - ESGF search API for bulk ingress
- Synda tool for bulk and parallel ingress (documentation can be found [HERE](#))


We're gonna need a bigger bucket!

Next Steps: Ingress - Expanding our Bucket



- Proposal submitted and pending for adding another 3 PB for esgf-world S3 bucket and Zarr S3 bucket for a total of 5 PB
- Comparison tracker of the two buckets can be found [HERE](#)


ESGF node: Future Architecture



Community-driven research efforts: Bring analysis to data



ESGF federated nodes across the world



<https://earthdata.nasa.gov/>

Towards a more accessible, discoverable and performant CMIP6 data holding in the cloud.

Earth System Grid Federation (ESGF) in the cloud. -By Kristopher Rand

Towards a generic, flexible API to create analysis-ready cloud-optimized dataset. - By Charles Stern

Speed up your multi-model analysis, using CMIP6 pre-processor and Dask - By Julius Busecke